

Tree-based methods

Hastie, Tibshirani, Friedman Ch 9.2, Ch 10

Kevin Murphy Ch. 16.1-16.4

James, Witten, Hastie, Tibshirani Ch 8

CS 6140

Machine Learning

Professor Olga Vitek

March 23, 2017

Decision trees

See handout (no author)

Overview: recursive partitioning

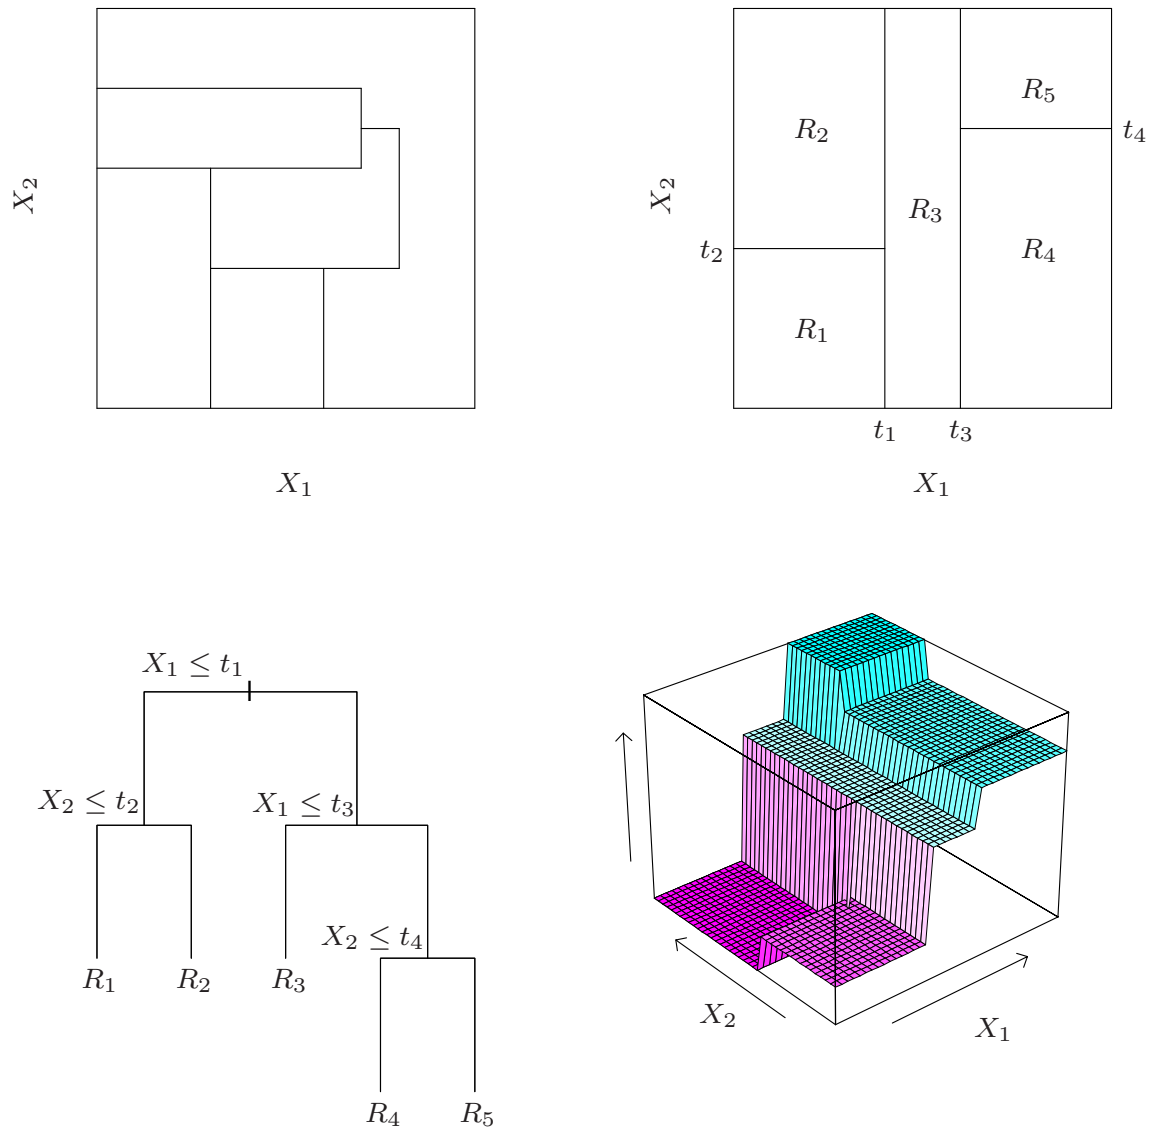


Fig. 9.2. Hastie, Tibshirani, Friedman
The Elements of Statistical Learning 2008

Example:

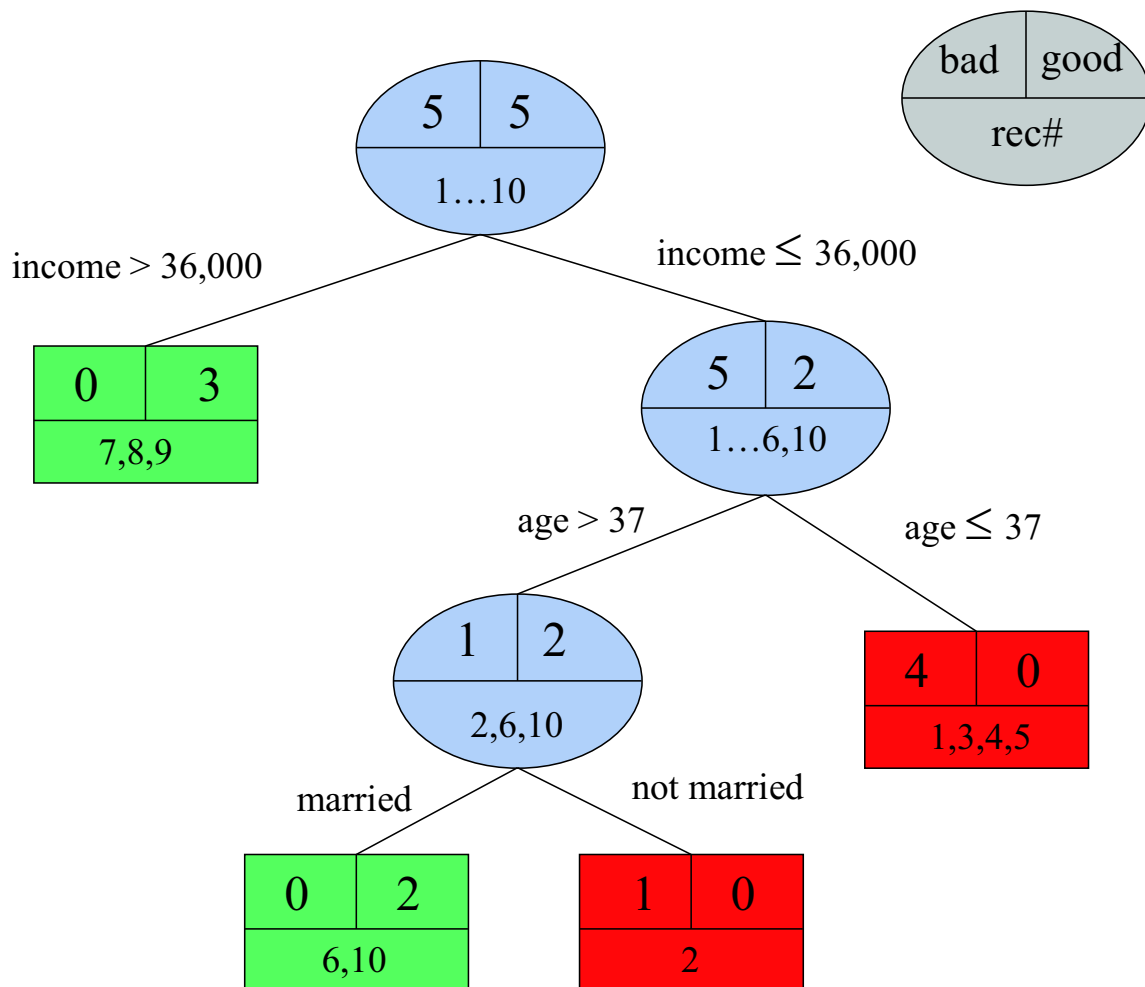
categorical response

credit score classification

Record	age	married?	own house	income	gender	class
1	22	no	no	28,000	male	bad
2	46	no	yes	32,000	female	bad
3	24	yes	yes	24,000	male	bad
4	25	no	no	27,000	male	bad
5	29	yes	yes	32,000	female	bad
6	45	yes	yes	30,000	female	good
7	63	yes	yes	58,000	male	good
8	36	yes	no	52,000	male	good
9	23	no	yes	40,000	female	good
10	50	yes	yes	28,000	female	good

Anonymous. See handout.

Credit score classification




Anonymous. See handout.

Quality of split

- Choose split to minimize node “impurity”
 - Define as function of relative class frequencies
 $i(t) = \phi(p_1, p_2, \dots, p_J)$ with J classes
 - $i(t)$ maximized at $(\frac{1}{J}, \frac{1}{J}, \dots, \frac{1}{J})$
 - $i(t)$ minimized at $(0, 0, \dots, 1)$ (for some class)
 - $i(t)$ is symmetric function of (p_1, p_2, \dots, p_J)
- Quality of split s at node t
 - $\Delta i(s, t) = i(t) - \pi(l) i(l) - \pi(r) i(r)$, where
 $\pi(l)$ is the proportion of points sent to the left
 $\pi(r)$ is the proportion of points sent to the right

Measures of node impurity

- Resubstitution error: $i(t) = 1 - \max_j p(j|t)$
 - $p(j|t)$: relative frequency of class j in node t
 - $i(t)$: % of misclassified cases
 - Node impurity simplifies to $\Delta i(s, t) = \max_j p(j|l) \pi(l) + \max_j p(j|r) \pi(r) - \max_j p(j|t)$
 - Problem: ignores where misclassification occurs
 - Below: same impurity; prefer split to the right
 - Ideally, ϕ would be concave (impurity would decrease faster than linearly)



Concave measures of impurity

- Gini index

- Two classes:

$$i(t) = p(0|t) p(1|t) = p(0|t) (1 - p(0|t))$$

- Multiple classes

$$i(t) = \sum_{j=1}^J p(j|t) (1 - p(j|t))$$

- Variance of Bernoulli drawing from this class

- Entropy

- Two classes:


$$i(t) = -p(0|t) \log p(0|t) - p(1|t) \log p(1|t)$$

- Multiple classes

$$i(t) = - \sum_{j=1}^J p(j|t) \log p(j|t)$$

- Average amount of info gathered by drawing a point from the node. Pure node \rightarrow no info.

Measures of node impurity



Two-class example; Entropy scaled to max=0.5

General properties:

- $\phi(0) = \phi(1) = 0$
- $\phi(p) = \phi(p(1 - p))$
- $\phi''(p) < 0, 0 < p < 1$

Fig. 9.3. Hastie, Tibshirani, Friedman

The Elements of Statistical Learning 2008

Search for splits

Income	Class	Quality (split after) 0.25–
24	B	$0.1(1)(0) + 0.9(4/9)(5/9) = 0.03$
27	B	$0.2(1)(0) + 0.8(3/8)(5/8) = 0.06$
28	B,G	$0.4(3/4)(1/4) + 0.6(2/6)(4/6) = 0.04$
30	G	$0.5(3/5)(2/5) + 0.5(2/5)(3/5) = 0.01$
32	B,B	$0.7(5/7)(2/7) + 0.3(0)(1) = 0.11$
40	G	$0.8(5/8)(3/8) + 0.2(0)(1) = 0.06$
52	G	$0.9(5/9)(4/9) + 0.1(0)(1) = 0.03$
58	G	

- Split on one predictor at a time
 - X numeric:
 - $X \leq \text{constant}$ for *constant* in range of X
 - In practice, split wrt distinct values of X
 - X categorical with values in V :
 - $X \in S$, S any subset of V
 - Number of possible splits is finite

Anonymous. See handout.

Tree construction

Algorithm: Construct tree

nodelist \leftarrow {training sample}

Repeat

 current node \leftarrow select node from nodelist

 nodelist \leftarrow nodelist $-$ current node

 if impurity(current node) > 0

 then

$S \leftarrow$ candidate splits in current node

$s^* \leftarrow \arg \max_{s \in S} \text{impurity reduction}(s, \text{current node})$

 child nodes \leftarrow apply(s^* , current node)

 nodelist \leftarrow nodelist \cup child nodes

 fi

Until nodelist = \emptyset

- Local greedy search; globally suboptimal.
- (Partial) remedy: grow maximal tree, then prune

P	Q	$P \text{ xor } Q$
1	1	0
1	0	1
0	1	1
0	0	0

Anonymous. See handout.

Cost-complexity pruning

- Notation
 - Tree T , and maximal tree T_{max}
 - $R(T)$ % of misclassified training set cases
 - $|T|$ tree size (i.e. # of terminal nodes)
 - Total cost of the tree $C_\alpha(T) = R(T) + \alpha|T|$
 - α : parameter penalizing tree complexity
- For every α , there exists a smallest subtree $T(\alpha)$ of T_{max} , such that:
 - No subtree of T_{max} has lower cost than $T(\alpha)$:
$$C_\alpha(T(\alpha)) = \min_{T \leq T_{max}} C_\alpha(T)$$
 - If there is a tie, we pick the smallest tree:
If $C_\alpha(T) = C_\alpha(T(\alpha))$, then $T(\alpha) \leq T$
- Consequence
 - It is impossible to have two non-nested subtrees with same min cost
 - Although α is continuous, only need a final set that changes tree structure
 $\{\alpha_1, \alpha_2, \dots\} \rightarrow T_1 > T_2 > \dots \{t_1\}$

Cost-complexity pruning

- Cost of T_t :
 - as an intermediate node: $C_\alpha(T_t) = R(T_t) + \alpha|T_t|$
 - as a terminal node: $C_\alpha(t) = R(t) + \alpha \cdot 1$
- The pruned tree has the same cost-complexity as the original tree when $C_\alpha(t) = C_\alpha(T_t)$

$$\begin{aligned} R(T_t) + \alpha|T_t| &= R(t) + \alpha \cdot 1 \\ \alpha &= \frac{R(t) - R(T_t)}{T - 1} \end{aligned}$$

- For any t , when we increase α beyond this level, pruned tree is better

Cost-complexity pruning

Algorithm: Compute tree sequence

$$T_1 \leftarrow T(0)$$

$$\alpha_1 \leftarrow 0$$

$$k \leftarrow 1$$

While $T_k > \{t_1\}$ do

For all non-terminal nodes $t \in T_k$

$$g_k(t) \leftarrow \frac{R(t) - R(T_{k,t})}{(|\tilde{T}_{k,t}| - 1)}$$

$$\alpha_{k+1} \leftarrow \min_t g_k(t)$$

Visit the nodes in top-down order and prune
whenever $g_k(t) = \alpha_{k+1}$ to obtain T_{k+1}

$$k \leftarrow k + 1$$

od

Example

