

Tree-based methods

Hastie, Tibshirani, Friedman Ch 9.2, Ch 10

Kevin Murphy Ch. 16.1-16.4

James, Witten, Hastie, Tibshirani Ch 8

CS 6140

Machine Learning

Professor Olga Vitek

April 6, 2017

Decision trees

See handout (no author)

Overview: recursive partitioning

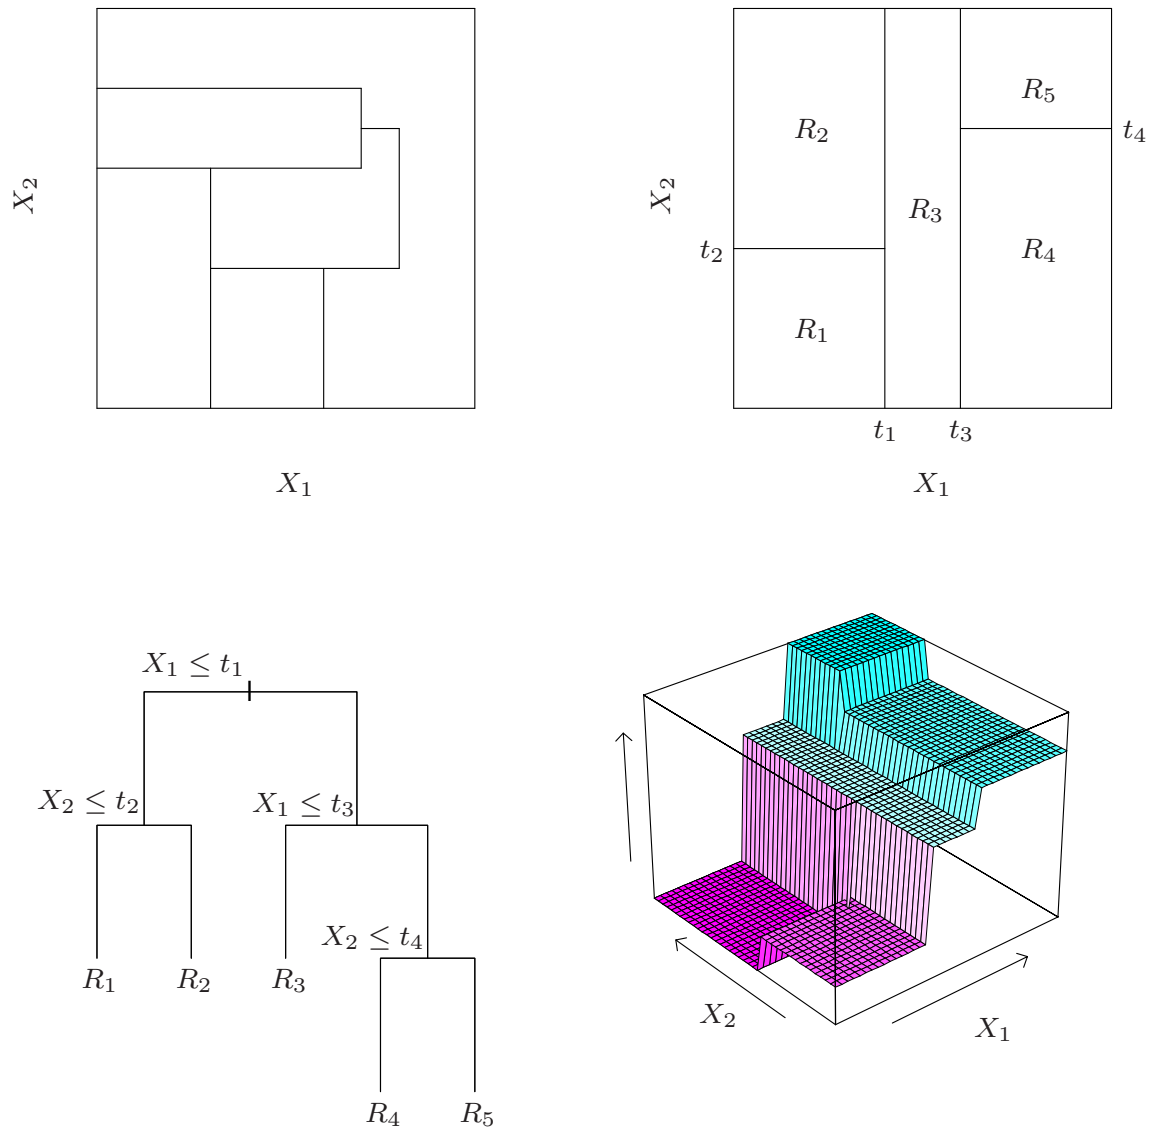


Fig. 9.2. Hastie, Tibshirani, Friedman
The Elements of Statistical Learning 2008

Example:

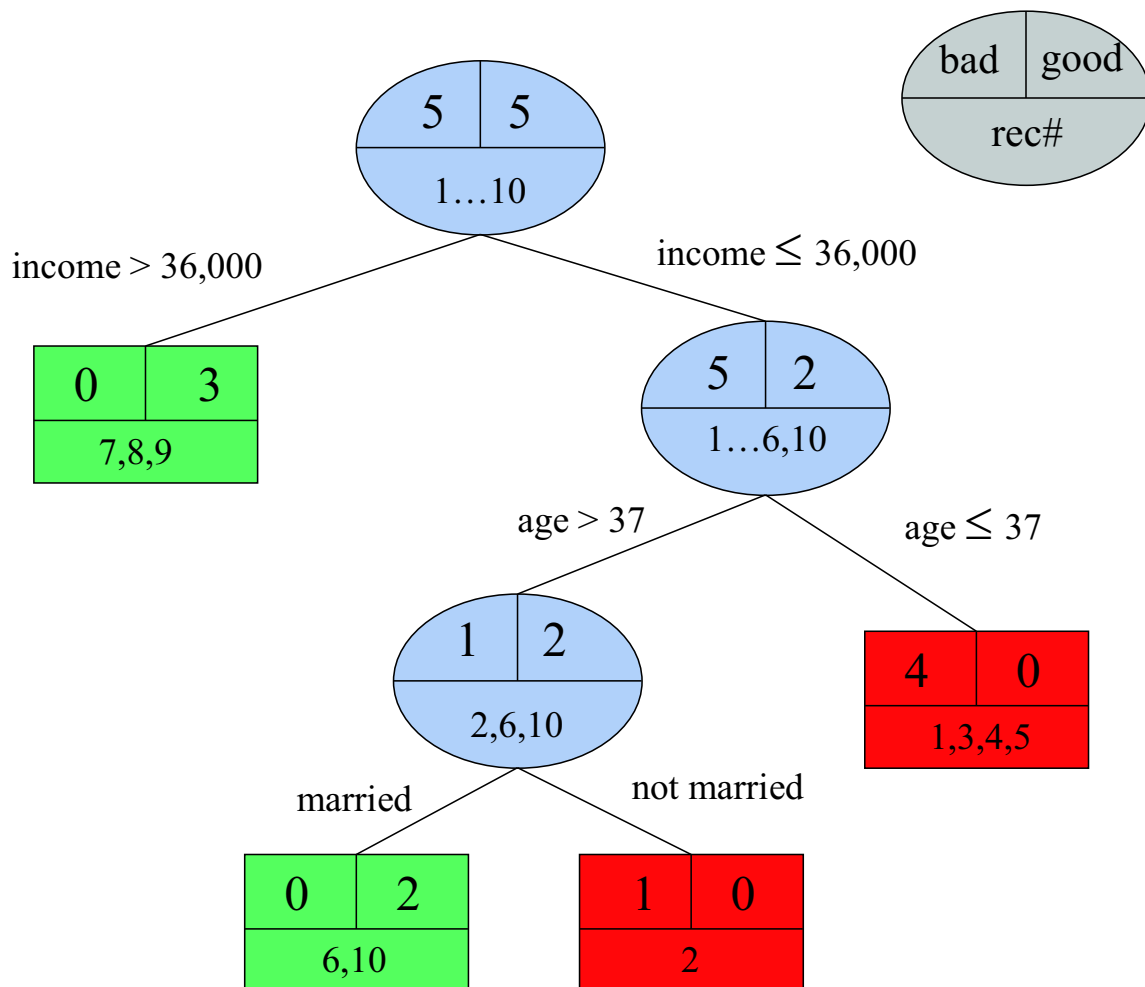
categorical response

credit score classification

Record	age	married?	own house	income	gender	class
1	22	no	no	28,000	male	bad
2	46	no	yes	32,000	female	bad
3	24	yes	yes	24,000	male	bad
4	25	no	no	27,000	male	bad
5	29	yes	yes	32,000	female	bad
6	45	yes	yes	30,000	female	good
7	63	yes	yes	58,000	male	good
8	36	yes	no	52,000	male	good
9	23	no	yes	40,000	female	good
10	50	yes	yes	28,000	female	good

Anonymous. See handout.

Credit score classification




Anonymous. See handout.

Quality of split

- Choose split to minimize node “impurity”
 - Define as function of relative class frequencies
 $i(t) = \phi(p_1, p_2, \dots, p_J)$ with J classes
 - $i(t)$ maximized at $(\frac{1}{J}, \frac{1}{J}, \dots, \frac{1}{J})$
 - $i(t)$ minimized at $(0, 0, \dots, 1)$ (for some class)
 - $i(t)$ is symmetric function of (p_1, p_2, \dots, p_J)
- Quality of split s at node t
 - $\Delta i(s, t) = i(t) - \pi(l) i(l) - \pi(r) i(r)$, where
 $\pi(l)$ is the proportion of points sent to the left
 $\pi(r)$ is the proportion of points sent to the right

Measures of node impurity

- Resubstitution error: $i(t) = 1 - \max_j p(j|t)$
 - $p(j|t)$: relative frequency of class j in node t
 - $i(t)$: % of misclassified cases
 - Node impurity simplifies to $\Delta i(s, t) = \max_j p(j|l) \pi(l) + \max_j p(j|r) \pi(r) - \max_j p(j|t)$
 - Problem: ignores where misclassification occurs
 - Below: same impurity; prefer split to the right
 - Ideally, ϕ would be concave (impurity would decrease faster than linearly)



Left: $\Delta i = 1 - \frac{1}{2} - \frac{1}{2} \left(1 - \frac{3}{4}\right) - \frac{1}{2} \left(1 - \frac{3}{4}\right) = 0.25$

Right: $\Delta i = 1 - \frac{1}{2} - \frac{3}{4} \left(1 - \frac{2}{3}\right) - \frac{1}{4} (1 - 1) = 0.25$

Concave impurity measures

- Gini index

- Two classes:

$$i(t) = p(0|t) \cdot p(1|t) = p(0|t) \cdot (1 - p(0|t))$$

- Multiple classes

$$i(t) = \sum_{j=1}^J p(j|t) \cdot (1 - p(j|t))$$

- Variance of Bernoulli drawing from this class
Impurity reduction \rightarrow variance reduction

- Entropy

- Two classes:

$$i(t) = -p(0|t) \cdot \log p(0|t) - p(1|t) \cdot \log p(1|t)$$


- Multiple classes

$$i(t) = -\sum_{j=1}^J p(j|t) \cdot \log p(j|t)$$

- Average amount of info gathered by drawing (with replacement) a point from the node and recording its class.
 \rightarrow Max impurity reduction = min loss of info.

- Low impurity only for very “pure” nodes

Measures of node impurity



Two-class example; Entropy scaled to max=0.5

General properties:

- $\phi(0) = \phi(1) = 0$
- $\phi(p) = \phi(p(1 - p))$
- $\phi''(p) < 0, 0 < p < 1$

Fig. 9.3. Hastie, Tibshirani, Friedman

The Elements of Statistical Learning 2008

Search for splits (Gini)

Income	Class	Quality (split after) 0.25–
24	B	$0.1(1)(0) + 0.9(4/9)(5/9) = 0.03$
27	B	$0.2(1)(0) + 0.8(3/8)(5/8) = 0.06$
28	B,G	$0.4(3/4)(1/4) + 0.6(2/6)(4/6) = 0.04$
30	G	$0.5(3/5)(2/5) + 0.5(2/5)(3/5) = 0.01$
32	B,B	$0.7(5/7)(2/7) + 0.3(0)(1) = 0.11$
40	G	$0.8(5/8)(3/8) + 0.2(0)(1) = 0.06$
52	G	$0.9(5/9)(4/9) + 0.1(0)(1) = 0.03$
58	G	

- Split on one predictor at a time
 - X numeric:
 - $X \leq \text{constant}$ for *constant* in range of X
 - In practice, split wrt distinct values of X
 - X categorical with values in V :
 - $X \in S$, S any subset of V
 - Number of possible splits is finite

Anonymous. See handout.

Tree construction

Algorithm: Construct tree

nodelist \leftarrow {training sample}

Repeat

 current node \leftarrow select node from nodelist

 nodelist \leftarrow nodelist $-$ current node

 if impurity(current node) > 0

 then

$S \leftarrow$ candidate splits in current node

$s^* \leftarrow \arg \max_{s \in S} \text{impurity reduction}(s, \text{current node})$

 child nodes \leftarrow apply(s^* , current node)

 nodelist \leftarrow nodelist \cup child nodes

 fi

Until nodelist = \emptyset

- Local greedy search; globally suboptimal

Example: true labels $P \text{ xor } Q$ can't be classified by splitting on P or Q

- (Partial) remedy: grow maximal tree, then prune

P	Q	$P \text{ xor } Q$
1	1	0
1	0	1
0	1	1
0	0	0

Cost-complexity pruning

- Notation

- Tree T , and maximal tree T_{max}
- $R(T)$ resubstitution error on training set
- $|T|$ tree size (i.e. # of terminal nodes)
- Total cost of the tree $C_\alpha(T) = R(T) + \alpha|T|$
- α : parameter penalizing tree complexity

- For every α , there exists a smallest subtree $T(\alpha)$ of T_{max} , such that:

- No subtree of T_{max} has lower cost than $T(\alpha)$:
$$C_\alpha(T(\alpha)) = \min_{T \leq T_{max}} C_\alpha(T)$$
- If there is a tie, we pick the smallest tree:
If $C_\alpha(T) = C_\alpha(T(\alpha))$, then $T(\alpha) \leq T$

- Consequence

- It is impossible to have two non-nested subtrees with same min cost
- Now vary α . Although α is continuous, only need a final set that changes tree structure
 $\{\alpha_1, \alpha_2, \dots\} \rightarrow T_1 > T_2 > \dots \{t_1\}$

Cost-complexity pruning

- T_1 : min subtree of T_{max} with same resubstitution error
- Cost of T_t :
 - as an intermediate node: $C_\alpha(T_t) = R(T_t) + \alpha|T_t|$
 - as a terminal node: $C_\alpha(t) = R(t) + \alpha \cdot 1$
- The pruned tree has the same cost-complexity as the original tree when $C_\alpha(t) = C_\alpha(T_t)$

$$R(T_t) + \alpha|T_t| = R(t) + \alpha \cdot 1$$
$$\alpha = \frac{R(t) - R(T_t)}{|T| - 1}$$

- For any t , when we increase α beyond this level, pruned tree is better
 - Obtain the next tree by pruning the current
 - T_k is the smallest minimizing subtree for $\alpha \in [\alpha_k, \alpha_{k+1})$
- There are $\lfloor 1.5028369^{|T|} \rfloor$ pruned trees

Cost-complexity pruning

Algorithm: Compute T_1 from T_{max}

$T' \leftarrow T_{max}$

Repeat

Pick any pair of terminal nodes ℓ and r with common parent t in T'
such that $R(t) = R(\ell) + R(r)$, and set

$T' \leftarrow T' - T_t$ (i.e. prune T' in t)

Until no more such pair exists

$T_1 \leftarrow T'$

Algorithm: Compute tree sequence

$T_1 \leftarrow T(0)$

$\alpha_1 \leftarrow 0$

$k \leftarrow 1$

While $T_k > \{t_1\}$ do

For all non-terminal nodes $t \in T_k$

$$g_k(t) \leftarrow \frac{R(t) - R(T_{k,t})}{(|\tilde{T}_{k,t}| - 1)}$$

$\alpha_{k+1} \leftarrow \min_t g_k(t)$

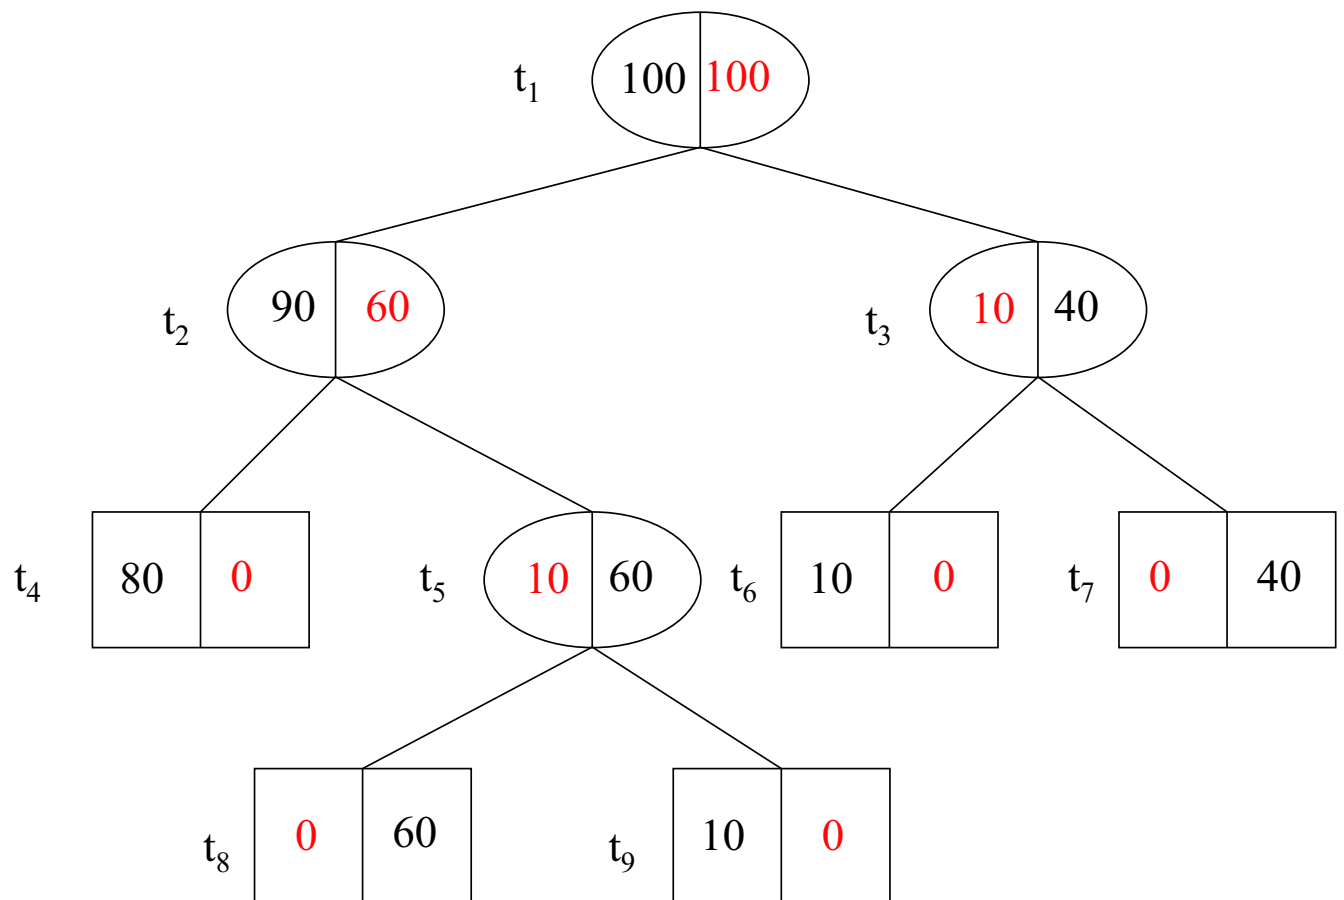
Visit the nodes in top-down order and prune

whenever $g_k(t) = \alpha_{k+1}$ to obtain T_{k+1}

$k \leftarrow k + 1$

od

Example: initial tree



Example: sequential pruning

- $k = 1$

$$g_1(t_5) = \frac{\frac{70}{200} \cdot \frac{10}{70} - \frac{70}{200} \cdot \frac{60}{70} \cdot 0 - \frac{70}{200} \cdot \frac{10}{70} \cdot 0}{2 - 1} = \frac{1}{20}$$

$$g_1(t_3) = \frac{\frac{50}{200} \cdot \frac{10}{50} - \frac{50}{200} \cdot \frac{10}{50} \cdot 0 - \frac{50}{200} \cdot \frac{40}{50} \cdot 0}{2 - 1} = \frac{1}{20}$$

$$g_1(t_2) = \frac{\frac{150}{200} \cdot \frac{60}{150} - \frac{150}{200} \cdot \frac{80}{150} \cdot 0 - \frac{150}{200} \cdot \frac{70}{150} \cdot \frac{60}{70} \cdot 0 - \frac{150}{200} \cdot \frac{70}{150} \cdot \frac{10}{70} \cdot 0}{3 - 1}$$

$$= \frac{3}{20}$$

\Rightarrow prune t_3 and t_5

- $k = 2$

$$g_2(t_2) = \frac{\frac{150}{200} \cdot \frac{60}{150} - \frac{150}{200} \cdot \frac{80}{150} \cdot 0 - \frac{150}{200} \cdot \frac{70}{150} \cdot \frac{10}{70}}{2 - 1} = \frac{5}{20}$$

$$g_2(t_1) = \frac{1 \cdot \frac{100}{200} - \frac{80}{200} \cdot 0 - \frac{70}{200} \cdot \frac{10}{70} - \frac{50}{200} \cdot \frac{10}{50}}{3 - 1} = \frac{4}{20}$$

\Rightarrow prune t_1

- $\alpha: \{ \alpha_1 = 0, \alpha_2 = \frac{1}{20}, \alpha_3 = \frac{4}{20} \}$

Selecting optimal tree

Algorithm:

- Construct a full tree on the full dataset
- Compute $\{\alpha_1, \alpha_2, \dots\}$
- Obtain $\{T_1, T_2, \dots, t\}$ for each α
- Define sequence $\{\beta_j = \sqrt{\alpha_{j-1} \cdot \alpha_j}\}$
- Divide data into folds
- Within each fold:
 - Build a sequence of trees over β_j
 - Compute prediction error for left-out observations
- For each β_j , sum $\#$ misclassifications over folds
- Select β_j with min $\#$ mis-classifications
- Report the β_j sub-tree of the full data

Challenges: Unstable (i.e., high variance), not smooth, can't capture additive structures

Bagging (=bootstrap aggregation)

- Motivation

- Average prediction on B independent datasets
↓ prediction variance:

$$\hat{f}_{avg} = \frac{1}{B} \sum_{b=1}^B \hat{f}^b(x), \quad Var\{\hat{f}_{avg}\} = \frac{Var\{\hat{f}^b(x)\}}{B}$$


- Bootstrap (HTF Sec 8.7)

- Mimics by sampling with replacement B datasets from the original dataset to ↓ prediction variance
- Grows maximal trees

- Features

- + *Out-of-bag error*: in each iteration b , use left-out observations to estimate prediction error (for each observation, average the error over all times when it was left out)
- + *Variable importance*: total amount of decrease in impurity or RSS by splitting on the predictor (averaged over trees)
- *Loss of interpretation*: a bagged tree \neq a tree

Bootstrap



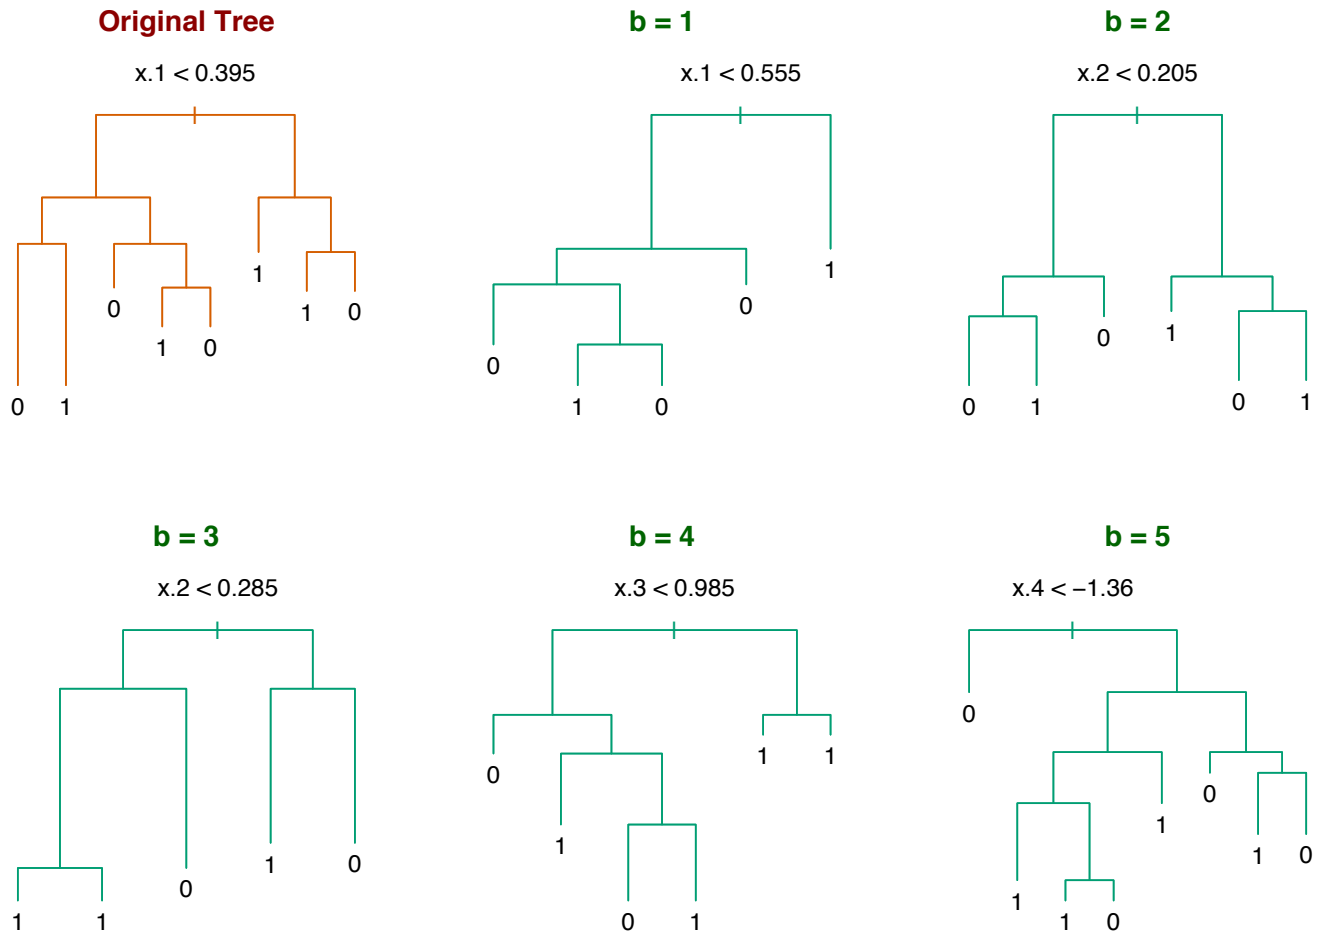
$$\hat{f}_{bag} = \frac{1}{B} \sum_{b=1}^B \hat{f}^{*b}(x)$$

Under square loss, for a classifier $\hat{f}^*(x)$ (HTF Sec 8.7.1):

$$\begin{aligned} E\{Y - \hat{f}^*(x)\}^2 &= E\{Y - \hat{f}_{bag}(x) + \hat{f}_{bag}(x) - \hat{f}^*(x)\}^2 \\ &= E\{Y - \hat{f}_{bag}(x)\}^2 + E\{\hat{f}_{bag}(x) - \hat{f}^*(x)\}^2 \\ &\geq E\{Y - \hat{f}_{bag}(x)\}^2 \end{aligned}$$

Does not always hold for 0-1 cost (misclassification)

Simulation




$N = 30$; up to 200 bootstrap samples

True model:

$$P\{Y = 1 | x_1 \leq 0.5\} = 0.2 \text{ and } P\{Y = 1 | x_1 > 0.5\} = 0.8$$

HTF Fig. 8.9

Simulation



$N = 30$; up to 200 bootstrap samples


True model:

$$P\{Y = 1|x_1 \leq 0.5\} = 0.2 \text{ and } P\{Y = 1|x_1 > 0.5\} = 0.8$$

Random forest


- Motivation
 - In bagging, prediction from trees can be *correlated* in presence of single dominant predictor
 - Ave. correlated predictions does not \downarrow variance
- Random forest
 - Introduces diversity in trees
 - Randomly selects $m \approx \sqrt{p}$ predictors per split
 - On average $(p-m)/p$ splits skip strong predictor
 - Random forest with $(m = p) =$ bagging

Example (experimental data)



Dashed line: single tree. OOB is too optimistic here.
JWHT Fig.8.8

Example (experimental data)




5-class gene expression data set with $p = 500$ predictors. Random forest ($m < p$) slightly improved over bagging ($m = p$). Single tree error rate: 45.7 %.

JWHT Fig.8.10

Motivation of shrinkage: James-Stein estimator

- Motivation:
 - Common wisdom: the best guess about the future is average of the past
 - Stein's paradox: average is sub-optimal in multivariate settings
- Set-up: I random variables
 - $Y_1 \stackrel{iid}{\sim} \mathcal{N}(\mu_1, \sigma^2), \dots, Y_I \stackrel{iid}{\sim} \mathcal{N}(\mu_I, \sigma^2)$
 - Goal: jointly estimate μ_1, \dots, μ_I with $\bar{Y}_1, \dots, \bar{Y}_I$.
 - Result: $\text{MSE} = E\{\|\hat{\mu} - \mu\|^2\}$ is not minimized
- Solution: can reduce MSE by
 - $\hat{\mu}_i = (1 - c) \cdot \bar{\mu} + c \cdot \hat{\mu}_i = \bar{\mu} + c \cdot (\hat{\mu}_i - \bar{\mu})$
 - weighted ave. of i th mean and overall mean
 - Here $c = 1 - \frac{(I-3) \cdot \sigma^2}{\sum_{i=1}^I (\hat{\mu}_i - \bar{\mu})^2}$
 - For fixed I and σ^2 , $c \uparrow$ if between-means variance $(\hat{\mu}_i - \bar{\mu})^2 \uparrow$
- c is the “shrinking” factor

Motivation of shrinkage: James-Stein estimator



$\bar{y} = 0.265; c = 0.212$
The shrinkage is $\approx 20\%$


The paper shows that shrunken estimates are closer to the “true” seasonal averages than the average-based estimators.

Efron & Morris, 1977

Extensions of shrinkage

- Shrinking to other target values
 - Any target value besides $\bar{\mu}$, e.g. to 0: $\hat{\mu}_i = c \cdot \hat{\mu}_i$
- Observations with unequal variance
 - $Y_1 \stackrel{iid}{\sim} \mathcal{N}(\mu_1, \sigma_1^2), \dots, Y_I \stackrel{iid}{\sim} \mathcal{N}(\mu_I, \sigma_I^2)$
 - c inversely related to the variance of the variable
 - More uncertainty \rightarrow more shrinkage
 - If a value is unusually large, it can be more reasonably attributed to random variation than to value of the mean
- Recently extended to non-Normal r.v
- Reduces the overall risk
 - Introduces error when one mean is atypical
- Is equivalent to specifying a prior distribution on the overall mean
 - Sometimes called “empirical Bayes”

Extensions of shrinkage: James-Stein estimator



Shrinking of random variables with unequal variance

Efron & Morris, 1977

Boosting

- Motivation
 - Trees, bagging, RF view obs. as exchangeable
 - Can overlook a few hard-to-classify observations
- Boosting
 - Sequentially grow trees
 - At each step, \uparrow weight of unexplained obs.
 - “Learn slowly”
 - Each next tree depends on the previous tree
- Parameters:
 - Number of trees B
Large B tends to overfit
 - Weight parameter λ
 λ can be viewed as a shrinkage parameter
 - Number of splits d (interaction depth).
Tree with $d = 1 =$ stump
If $d = 1$, each tree involves one variable
→ additive model
- Generalization:
 - Same approach can be applied to other modeling strategies

Boosting algorithm

Algorithm 8.2 *Boosting for Regression Trees*

1. Set $\hat{f}(x) = 0$ and $r_i = y_i$ for all i in the training set.
2. For $b = 1, 2, \dots, B$, repeat:
 - (a) Fit a tree \hat{f}^b with d splits ($d + 1$ terminal nodes) to the training data (X, r) .
 - (b) Update \hat{f} by adding in a shrunk version of the new tree:

$$\hat{f}(x) \leftarrow \hat{f}(x) + \lambda \hat{f}^b(x). \quad (8.10)$$

- (c) Update the residuals,


$$r_i \leftarrow r_i - \lambda \hat{f}^b(x_i). \quad (8.11)$$

3. Output the boosted model,

$$\hat{f}(x) = \sum_{b=1}^B \lambda \hat{f}^b(x). \quad (8.12)$$

JWHT Algorithm 8.2

Example (experimental data)




Boosting outperformed Random Forest. However, SE (not shown) is large enough to make the methods comparable


JWHT Fig.8.11

Example (simulated data)

Bagged Decision Rule



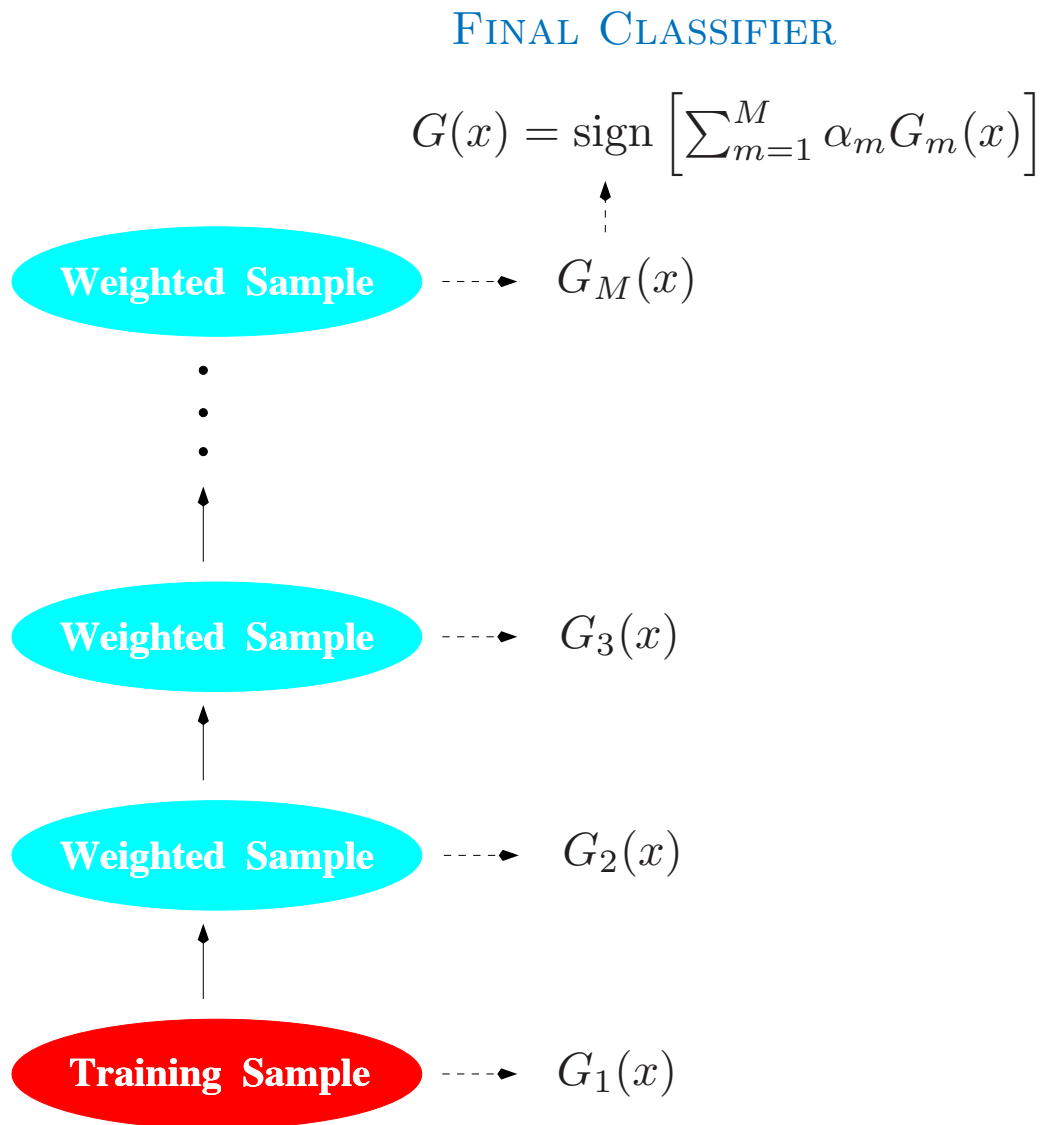
Boosted Decision Rule



Bagging 0-1 decision rule (i.e., misclassification rate) over $B = 50$ bootstrap samples. The trees split around the middle range of X_1 or X_2 , and have little contribution away from the center.

HTF Fig.8.12

General approach: AdaBoost



$Y \in \{-1, 1\}$, classifier $G(X)$, $\overline{\text{err}} = \frac{1}{N} \sum_{i=1}^N I\{y_i \neq G(x_i)\}$

Train many weak classifiers \rightarrow use a “committee” vote

HTF Chapter 10; Fig. 10.2

More generally: Forward stagewise additive modeling

- Basis function expansion

$$f(x) = \sum_{m=1}^M \beta_m b(x, \gamma_m)$$

- β_m : expansion coefficients
- $b(x, \gamma_m)$ simple basis functions,
e.g. classifier $G_n \in \{-1, 1\}$

- Loss function

- Generally, optimize over all parameters (β_m, γ_m)
of all basis functions

$$\min_{\beta_m, \gamma_m} \sum_{i=1}^N L \left(y_i, \sum_{m=1}^M \beta_m b(x_i, \gamma_m) \right)$$

- Often simplified: optimizing one basis function

$$\min_{\beta, \gamma} \sum_{i=1}^N L \left(y_i, \sum_{m=1}^M \beta b(x_i, \gamma) \right)$$

Forward stagewise additive modeling

Algorithm 10.2 *Forward Stagewise Additive Modeling.*

1. Initialize $f_0(x) = 0$.

2. For $m = 1$ to M :

(a) Compute

$$(\beta_m, \gamma_m) = \arg \min_{\beta, \gamma} \sum_{i=1}^N L(y_i, f_{m-1}(x_i) + \beta b(x_i; \gamma)).$$

(b) Set $f_m(x) = f_{m-1}(x) + \beta_m b(x; \gamma_m)$.

For squared-error loss $L(y; f(x)) = (y - f(x))^2$:

$$\begin{aligned} L(y_i; f_{m-1}(x_i) + \beta b(x_i, \gamma)) &= (y_i - f_{m-1}(x_i) - \beta b(x_i, \gamma))^2 \\ &= (r_{im} - \beta b(x_i, \gamma))^2 \end{aligned}$$

$r_{im} = y_i - f_{m-1}(x_i)$: residual of m th model for i th obs.

$\beta b(x_i, \gamma)$: fits to the current residuals

Problem: squared-loss is bad for classification

HTF Chapter 10; Algorithm 10.2



Classification: AdaBoost

Algorithm 10.1 *AdaBoost.M1*.

1. Initialize the observation weights $w_i = 1/N$, $i = 1, 2, \dots, N$.
 2. For $m = 1$ to M :
 - (a) Fit a classifier $G_m(x)$ to the training data using weights w_i .
 - (b) Compute
$$\text{err}_m = \frac{\sum_{i=1}^N w_i I(y_i \neq G_m(x_i))}{\sum_{i=1}^N w_i}.$$
 - (c) Compute $\alpha_m = \log((1 - \text{err}_m)/\text{err}_m)$.
 - (d) Set $w_i \leftarrow w_i \cdot \exp[\alpha_m \cdot I(y_i \neq G_m(x_i))]$, $i = 1, 2, \dots, N$.
 3. Output $G(x) = \text{sign} \left[\sum_{m=1}^M \alpha_m G_m(x) \right]$.
-

- First, all observations have same weights
- At each step, \uparrow weights of misclassified observations by factor $\exp(\alpha_m)$
- It is a basis function expansion.
Basis functions: individual classifiers $G_m \in \{-1, 1\}$
- Minimize loss function, averaged over training data
- New aspect: w_i

AdaBoost weights



Exponential loss

- AdaBoost: equivalent to forward stagewise modeling with exponential loss $L(y; f(x)) = \exp\{-y f(x)\}$
- Optimization problem:
 $Y \in \{-1, 1\}$. For classifier $G_n \in \{-1, 1\}$, solve

$$\begin{aligned}(\beta_m, G_m) &= \min_{\beta, G} \sum_{i=1}^N \exp[-y_i (f_{m-1}(x_i) + \beta G(x))] \\&= \min_{\beta, G} \sum_{i=1}^N \exp[-y_i f_{m-1}(x_i)] \cdot \exp[-\beta y_i G(x)] \\&= \min_{\beta, G} \sum_{i=1}^N w_i^{(m)} \cdot \exp[-\beta y_i G(x)] \\&\quad \text{and } w_i^{(m)} = \exp[-y_i f_{m-1}(x_i)]\end{aligned}$$

Weights change with m , do not depend on β or G


- Solution

$$G_m = \arg \min_G \sum_{i=1}^N w_i^{(m)} I(y_i \neq G(x_i))$$

Same weights - see HTF Sec. 10.4

- Different loss funct. \rightarrow different properties

Alternative loss functions



Exponential loss:
attractive computational and analytical properties

$y \cdot f(x)$ is the “margin” in SVM

HTF Fig 10.4